

LogPulse AI - Incident Report

Generated : 2026-06-01 13:14:27 UTC

Trace ID : 1bd7618e-c07e-473c-adf9-4a790ac8fab8

1 - Incident Summary

Incident ID : 1bd7618e
Category : Memory
Severity : SeverityLevel.FATAL
Source : Kubernetes
Cache : HIT

2 - Raw Log

2026-01-15T22:14:07Z FATAL kubelet: Pod worker-api-7d9f8b-xk2q OOMKilled. Container exceeded memory limit of 256Mi. Exit code 137. Restart count: 14. Node: worker-1. Namespace: production.

3 - Classification

Category : Memory
Severity : SeverityLevel.FATAL
Summary : Kubernetes pod killed due to OOMKilled container exceeding its memory limit.
Source : Kubernetes

4 - Diagnosis

Root Cause : The container's configured memory limit of 256Mi is insufficient for its operational memory requirements, causing Kubernetes to repeatedly OOMKill the pod.

Confidence : 95%

Reasoning : The raw log and classification explicitly state the pod was OOMKilled for exceeding its 256Mi memory limit, which is directly addressed by the retrieved playbook titled 'Pod OOMKilled: Insufficient Memory Limit' and its recommended resolution of increasing the memory limit.

Source Tier : llm
Used History : No
Used Playbook : Yes

5 - Proposed Solution

Explanation:

The container's configured memory limit of 256Mi is insufficient for its operational memory requirements, causing

Kubernetes to repeatedly OOMKill the pod.

Steps:

- 1. Retrieve the detailed configuration and status of the OOMKilled pod to verify the exit code and termination reason.
- 2. Fetch the logs of the previously terminated container instance to analyze application behavior prior to the OOM event.
- 3. Check the current CPU and memory utilization of the pod.
- 4. Increase the memory limit of the worker-api deployment to 512Mi to prevent further OOM kills.

Commands:

```
$ kubectl get pod worker-api-7d9f8b-xk2q -n production -o yaml
$ kubectl logs worker-api-7d9f8b-xk2q -n production --previous --tail=100
$ kubectl top pod worker-api-7d9f8b-xk2q -n production
$ kubectl set resources deployment/worker-api -n production --limits=memory=512Mi
--requests=memory=256Mi
```

6 - Risk Assessment

Risk Level : RISKLEVEL.SAFE

Safe Commands (4):

- [OK] kubectl get pod worker-api-7d9f8b-xk2q -n production -o yaml
- [OK] kubectl logs worker-api-7d9f8b-xk2q -n production --previous --tail=100
- [OK] kubectl top pod worker-api-7d9f8b-xk2q -n production
- [OK] kubectl set resources deployment/worker-api -n production --limits=memory=512Mi --requests=memory=256Mi

Flagged Commands (0):

(none)

7 - Operator Decision

Operator rejected -- no execution.

8 - Historical Context

Times Seen : 2

First occurrence -- no prior incidents.

9 - Incident Timeline

```
[10:08:23] INCIDENT_RECEIVED -> Incident received (trace_id=1bd7618e...)
[10:08:23] CACHE_HIT -> Cache hit -- seen 2x
[10:08:28] RAG_RETRIEVED -> Playbook retrieved (filtered) -- Memory conf=0.82
[10:08:36] ROOT_CAUSE_GENERATED -> The container's configured memory limit of 256Mi is
insufficient for its operational memory requirements, causing Kubern
[10:09:27] SOLUTION_GENERATED -> 4 step(s), 4 command(s)
[10:09:27] RISK_ASSESSED -> safe -- 0 flagged
```

[10:14:27] WORKFLOW_COMPLETED -> Completed in 363877ms