

LogPulse AI - Incident Report

Generated : 2026-05-29 17:31:14 UTC

Trace ID : 2c78bd5a-eeed-4f7c-84e0-62b3e379b37f

1 - Incident Summary

Incident ID : 2c78bd5a
Category : Memory
Severity : SeverityLevel.FATAL
Source : Kubernetes
Cache : MISS

2 - Raw Log

17:30:21 FATAL OOMKilled: pod/worker-api exceeded memory limit 256Mi restarts=14

3 - Classification

Category : Memory
Severity : SeverityLevel.FATAL
Summary : Kubernetes worker API pod exceeded its memory limit and was terminated.
Source : Kubernetes

4 - Diagnosis

Root Cause : The pod/worker-api is repeatedly OOMKilled because its configured memory limit of 256Mi is insufficient for the application's runtime memory requirements.
Confidence : 99%
Reasoning : The raw log, classification, and incident history all confirm the pod is repeatedly OOMKilled due to exceeding its 256Mi memory limit, directly matching the previous root cause and the suggested playbook for insufficient memory limits.
Source Tier : llm
Used History : Yes
Used Playbook : Yes

5 - Proposed Solution

Explanation:

Gemini unavailable; remediation derived from playbook 'Pod OOMKilled: Insufficient Memory Limit' (RAG confidence=0.8040). Review the playbook steps and adapt commands to the live environment before execution.

Steps:

- 1. Verify the `OOMKilled` status using `kubectl get pod <pod-name> -o yaml`.
- 2. Check container memory usage metrics (e.g., Prometheus/Grafana) to identify peak memory consumption.
- 3. Analyze application logs for clues about memory spikes or growth patterns.
- 4. Increase the `resources.limits.memory` for the affected container in its Deployment/StatefulSet manifest.
- 5. Roll out the updated deployment using `kubectl apply -f <manifest-file>` or `kubectl rollout restart deployment/<deployment-name>`.
- 6. name: my-app

Commands:

```
$ kubectl get pod worker-api-7d9f8b-xk2q -o yaml
$ kubectl apply -f <manifest-file>
$ kubectl rollout restart deployment/worker-api
```

6 - Risk Assessment

Risk Level : **RISKLEVEL.WARNING**

Safe Commands (2):

```
[OK] kubectl get pod worker-api-7d9f8b-xk2q -o yaml
[OK] kubectl apply -f <manifest-file>
```

Flagged Commands (1):

```
[!] kubectl rollout restart deployment/worker-api [kubectl\s+rollout\s+restart]
```

7 - Operator Decision

Decision : OPERATORDECISION.APPROVED
Executed At : 2026-05-29T14:31:14.431368+00:00

Sandbox Results:

Command : kubectl get pod worker-api-7d9f8b-xk2q -o yaml
Exit : 0
Time : 0ms

Output:

```
[auto-approve] simulated: kubectl get pod worker-api-7d9f8b-xk2q -o yaml
```

Command : kubectl apply -f <manifest-file>
Exit : 0
Time : 0ms

Output:

```
[auto-approve] simulated: kubectl apply -f <manifest-file>
```

Command : kubectl rollout restart deployment/worker-api
Exit : 0
Time : 0ms

Output:

```
[auto-approve] simulated: kubectl rollout restart deployment/worker-api
```

8 - Historical Context

Times Seen : 1

Past Outcomes:

- analyzed on 2026-05-29 14:11:24.999019+00:00

9 - Incident Timeline

[14:31:14] WORKFLOW_COMPLETED -> Completed in 28850ms