

LogPulse AI - Incident Report

Generated : 2026-06-01 00:57:28 UTC

Trace ID : 9aecc65d-5a3e-4019-93c4-cf1e83a1f312

1 - Incident Summary

Incident ID : 9aecc65d
Category : Unknown
Severity : SeverityLevel.INFO
Source : Fallback
Cache : MISS

2 - Raw Log

2026-05-30 ERROR kubelet: Pod worker-api OOMKilled exceeded 256Mi limit

3 - Classification

Category : Unknown
Severity : SeverityLevel.INFO
Summary : Fallback classification triggered due to inference failure.
Source : Fallback

4 - Diagnosis

Root Cause : The pod 'worker-api' was OOMKilled because its configured memory limit of 256Mi was insufficient for its operational requirements.
Confidence : 90%
Reasoning : The raw log explicitly states 'Pod worker-api OOMKilled exceeded 256Mi limit', and the retrieved playbook, with high confidence, directly addresses 'Pods OOMKilled due to insufficient memory', recommending an increase in memory limits.
Source Tier : llm
Used History : No
Used Playbook : Yes

5 - Proposed Solution

Explanation:

The pod worker-api-7d9f8b-xk2q was OOMKilled because its configured memory limit of 256Mi was insufficient for its operational requirements.

Steps:

1. Inspect the pod details, termination reason, and resource limits
2. Retrieve logs from the container prior to the OOM crash to diagnose memory usage
3. Increase the memory limit to 512Mi and request to 256Mi to prevent OOM kills
4. Monitor the rolling update status of the worker-api deployment

Commands:

```
$ kubectl describe pod worker-api-7d9f8b-xk2q -n production
$ kubectl logs worker-api-7d9f8b-xk2q -n production --previous
$ kubectl set resources deployment/worker-api -n production --limits=memory=512Mi
--requests=memory=256Mi
$ kubectl rollout status deployment/worker-api -n production
```

6 - Risk Assessment

Risk Level : **RISKLEVEL.SAFE**

Safe Commands (4):

```
[OK] kubectl describe pod worker-api-7d9f8b-xk2q -n production
[OK] kubectl logs worker-api-7d9f8b-xk2q -n production --previous
[OK] kubectl set resources deployment/worker-api -n production --limits=memory=512Mi --requests=memory=256Mi
[OK] kubectl rollout status deployment/worker-api -n production
```

Flagged Commands (0):

(none)

7 - Operator Decision

Operator rejected -- no execution.

8 - Historical Context

Times Seen : 1

First occurrence -- no prior incidents.

9 - Incident Timeline

```
[21:57:05] INCIDENT_RECEIVED -> Incident received (trace_id=9aecc65d...)
[21:57:05] CACHE_MISS -> Cache miss -- routing to classifier
[21:57:06] INCIDENT_CLASSIFIED -> Fallback / Info
[21:57:08] RAG_RETRIEVED -> Playbook retrieved (fallback) conf=0.77
[21:57:13] ROOT_CAUSE_GENERATED -> The pod 'worker-api' was OOMKilled because its configured
memory limit of 256Mi was insufficient for its operational req
[21:57:27] SOLUTION_GENERATED -> 4 step(s), 4 command(s)
[21:57:27] RISK_ASSESSED -> safe -- 0 flagged
[21:57:27] WORKFLOW_COMPLETED -> Completed in 21861ms
```