

LogPulse AI - Incident Report

Generated : 2026-06-01 16:16:02 UTC

Trace ID : 64d92ac5-021e-40ac-ac45-89dce56095f8

1 - Incident Summary

Incident ID : 64d92ac5
Category : Unknown
Severity : SeverityLevel.INFO
Source : Fallback
Cache : MISS

2 - Raw Log

```
16:13:33 FATAL OOMKilled: pod/worker-api exceeded memory limit 256Mi restarts=14
```

3 - Classification

Category : Unknown
Severity : SeverityLevel.INFO
Summary : Fallback classification triggered due to inference failure.
Source : Fallback

4 - Diagnosis

Root Cause : The worker-api pod was configured with an insufficient memory limit of 256Mi, causing the Kubernetes OOM killer to repeatedly terminate the container as its memory usage exceeded this threshold.

Confidence : 95%

Reasoning : The raw log explicitly shows a FATAL OOMKilled event for pod/worker-api due to exceeding its 256Mi memory limit, leading to 14 restarts.

Source Tier : llm

Used History : No

Used Playbook : Yes

5 - Proposed Solution

Explanation:

The worker-api pod was configured with an insufficient memory limit of 256Mi, causing the Kubernetes OOM killer to repeatedly terminate the container as its memory usage exceeded this threshold.

Steps:

1. Retrieves the YAML manifest of the worker-api pod to verify its status and events, including OOMKilled.
2. Shows current resource usage (CPU/memory) for the worker-api pod and its containers.
3. Retrieves the last 100 lines of logs from the worker-api pod to look for memory-related events or errors.
4. Patches the worker-api deployment to increase the memory limit to 1Gi and request to 512Mi for its first container, triggering a rollout.
5. Restarts the worker-api deployment, forcing new pods to be created with the updated configuration.

Commands:

```
$ kubectl get pod worker-api-7d9f8b-xk2q -n production -o yaml
$ kubectl top pod worker-api-7d9f8b-xk2q -n production --containers
$ kubectl logs worker-api-7d9f8b-xk2q -n production --tail=100
$ kubectl patch deployment worker-api -n production --type='json' -p='[{"op": "replace", "path":
"/spec/template/spec/containers/0/resources/limits/memory", "value": "1Gi"}, {"op": "replace",
"path": "/spec/template/spec/containers/0/resources/requests/memory", "value": "512Mi"}]'
$ kubectl rollout restart deployment/worker-api -n production
```

6 - Risk Assessment

Risk Level : **RISKLEVEL.WARNING**

Safe Commands (4):

```
[OK] kubectl get pod worker-api-7d9f8b-xk2q -n production -o yaml
[OK] kubectl top pod worker-api-7d9f8b-xk2q -n production --containers
[OK] kubectl logs worker-api-7d9f8b-xk2q -n production --tail=100
[OK] kubectl patch deployment worker-api -n production --type='json' -p='[{"op": "replace", "path":
"/spec/template/spec/containers/0/resources/limits/memory", "value": "1Gi"}, {"op": "replace", "path":
"/spec/template/spec/containers/0/resources/requests/memory", "value": "512Mi"}]'
```

Flagged Commands (1):

```
[!] kubectl rollout restart deployment/worker-api -n production [kubectl\s+rollout\s+restart]
```

7 - Operator Decision

Operator rejected -- no execution.

8 - Historical Context

Times Seen : 1

First occurrence -- no prior incidents.

9 - Incident Timeline

```
[13:14:14] INCIDENT_RECEIVED -> Incident received (trace_id=64d92ac5...)
[13:14:14] CACHE_MISS -> Cache miss -- routing to classifier
[13:14:16] INCIDENT_CLASSIFIED -> Fallback / Info
[13:14:25] RAG_RETRIEVED -> Playbook retrieved (fallback) conf=0.78
```

[13:14:35] ROOT_CAUSE_GENERATED -> The worker-api pod was configured with an insufficient memory limit of 256Mi, causing the Kubernetes OOM killer to repea
[13:15:57] SOLUTION_GENERATED -> 5 step(s), 5 command(s)
[13:15:57] RISK_ASSESSED -> warning -- 1 flagged
[13:16:02] WORKFLOW_COMPLETED -> Completed in 108026ms