

LogPulse AI - Incident Report

Generated : 2026-06-01 14:41:49 UTC

Trace ID : 204f87ee-c3ab-465d-85c1-ef783013e04b

1 - Incident Summary

Incident ID : 204f87ee
Category : Unknown
Severity : SeverityLevel.INFO
Source : Fallback
Cache : MISS

2 - Raw Log

14:40:18 FATAL OOMKilled: pod/worker-api exceeded memory limit 256Mi restarts=14

3 - Classification

Category : Unknown
Severity : SeverityLevel.INFO
Summary : Fallback classification triggered due to inference failure.
Source : Fallback

4 - Diagnosis

Root Cause : The 'worker-api' pod is consistently exceeding its configured memory limit of 256Mi, leading to repeated OOMKills and restarts, indicating the application requires more memory than currently allocated.

Confidence : 95%

Reasoning : The raw log explicitly states 'OOMKilled' because the pod 'exceeded memory limit 256Mi' with 14 restarts. The retrieved playbook, with high confidence, directly addresses 'Pod OOMKilled: Insufficient Memory Limit' and suggests increasing memory limits as a resolution.

Source Tier : llm
Used History : No
Used Playbook : Yes

5 - Proposed Solution

Explanation:

6 - Risk Assessment

Risk Level : **RISKLEVEL.SAFE**

Safe Commands (0):

(none)

Flagged Commands (0):

(none)

7 - Operator Decision

Operator rejected -- no execution.

8 - Historical Context

Times Seen : 1

First occurrence -- no prior incidents.

9 - Incident Timeline

```
[11:40:42] INCIDENT_RECEIVED -> Incident received (trace_id=204f87ee...)
[11:40:42] CACHE_MISS -> Cache miss -- routing to classifier
[11:40:44] INCIDENT_CLASSIFIED -> Fallback / Info
[11:40:48] RAG_RETRIEVED -> Playbook retrieved (fallback) conf=0.78
[11:40:52] ROOT_CAUSE_GENERATED -> The 'worker-api' pod is consistently exceeding its configured
memory limit of 256Mi, leading to repeated OOMKills and re
[11:41:42] SOLUTION_GENERATED -> 0 step(s), 0 command(s)
[11:41:42] RISK_ASSESSED -> safe -- 0 flagged
[11:41:48] WORKFLOW_COMPLETED -> Completed in 66400ms
```