

LogPulse AI - Incident Report

Generated : 2026-06-01 17:02:23 UTC

Trace ID : a2fbecfb-e030-44c7-adb9-fc629b321e46

1 - Incident Summary

Incident ID : a2fbecfb
Category : Memory
Severity : SeverityLevel.FATAL
Source : Kubernetes
Cache : MISS

2 - Raw Log

16:55:45 FATAL OOMKilled: pod/worker-api exceeded memory limit 256Mi restarts=14

3 - Classification

Category : Memory
Severity : SeverityLevel.FATAL
Summary : Kubernetes worker API pod exceeded its memory limit and was terminated.
Source : Kubernetes

4 - Diagnosis

Root Cause : The worker-api pod is repeatedly exceeding its configured memory limit of 256Mi, leading to OOMKills and restarts, indicating the limit is insufficient for its operational needs.

Confidence : 95%

Reasoning : The raw log explicitly states 'OOMKilled: pod/worker-api exceeded memory limit 256Mi restarts=14'. The retrieved playbook is titled 'Pod OOMKilled: Insufficient Memory Limit' with high confidence (0.8153) and suggests increasing the memory limit. The 14 restarts confirm a recurring issue.

Source Tier : llm
Used History : No
Used Playbook : Yes

5 - Proposed Solution

Explanation:

6 - Risk Assessment

Risk Level : **RISKLEVEL.SAFE**

Safe Commands (0):

(none)

Flagged Commands (0):

(none)

7 - Operator Decision

Operator rejected -- no execution.

8 - Historical Context

Times Seen : 1

First occurrence -- no prior incidents.

9 - Incident Timeline

```
[13:55:47] INCIDENT_RECEIVED -> Incident received (trace_id=a2fbecfb...)
[13:55:47] CACHE_MISS -> Cache miss -- routing to classifier
[13:55:57] INCIDENT_CLASSIFIED -> Memory / Fatal
[13:55:59] RAG_RETRIEVED -> Playbook retrieved (filtered) -- Memory conf=0.82
[13:56:05] ROOT_CAUSE_GENERATED -> The worker-api pod is repeatedly exceeding its configured
memory limit of 256Mi, leading to OOMKills and restarts, indic
[13:57:23] SOLUTION_GENERATED -> 0 step(s), 0 command(s)
[13:57:23] RISK_ASSESSED -> safe -- 0 flagged
[14:02:23] WORKFLOW_COMPLETED -> Completed in 395758ms
```