

LogPulse AI - Incident Report

Generated : 2026-06-01 00:34:56 UTC

Trace ID : b8429e8a-59d2-4dac-b2a9-0bfb7321b1a8

1 - Incident Summary

Incident ID : b8429e8a
Category : Unknown
Severity : SeverityLevel.INFO
Source : Fallback
Cache : MISS

2 - Raw Log

2026-05-30 ERROR kubelet: Pod worker-api OOMKilled exceeded 256Mi limit

3 - Classification

Category : Unknown
Severity : SeverityLevel.INFO
Summary : Fallback classification triggered due to inference failure.
Source : Fallback

4 - Diagnosis

Root Cause : The worker-api pod was OOMKilled because its configured memory limit of 256Mi was insufficient for the application's operational requirements.
Confidence : 95%
Reasoning : The raw log explicitly states 'Pod worker-api OOMKilled exceeded 256Mi limit,' which directly matches the retrieved playbook 'Pods OOMKilled due to insufficient memory' with high confidence, whose primary resolution is to increase the memory limit.
Source Tier : llm
Used History : No
Used Playbook : Yes

5 - Proposed Solution

Explanation:

The worker-api pod was OOMKilled because its configured memory limit of 256Mi was insufficient for the application's operational requirements.

Steps:

1. Describe the affected pod to inspect its status, events, and termination reason.
2. Retrieve logs from the previous instance of the container before it was OOMKilled.
3. Increase the memory limit to 512Mi and request to 256Mi for the worker-api deployment to prevent OOMKills.
4. Monitor the rollout status of the updated worker-api deployment.

Commands:

```
$ kubectl describe pod worker-api-7d9f8b-xk2q -n production
$ kubectl logs worker-api-7d9f8b-xk2q -n production --previous
$ kubectl set resources deployment/worker-api -n production --limits=memory=512Mi
--requests=memory=256Mi
$ kubectl rollout status deployment/worker-api -n production
```

6 - Risk Assessment

Risk Level : **RISKLEVEL.SAFE**

Safe Commands (4):

```
[OK] kubectl describe pod worker-api-7d9f8b-xk2q -n production
[OK] kubectl logs worker-api-7d9f8b-xk2q -n production --previous
[OK] kubectl set resources deployment/worker-api -n production --limits=memory=512Mi --requests=memory=256Mi
[OK] kubectl rollout status deployment/worker-api -n production
```

Flagged Commands (0):

(none)

7 - Operator Decision

Operator rejected -- no execution.

8 - Historical Context

Times Seen : 1

First occurrence -- no prior incidents.

9 - Incident Timeline

```
[21:27:12] INCIDENT_RECEIVED -> Incident received (trace_id=b8429e8a...)
[21:27:12] CACHE_MISS -> Cache miss -- routing to classifier
[21:27:13] INCIDENT_CLASSIFIED -> Fallback / Info
[21:27:16] RAG_RETRIEVED -> Playbook retrieved (fallback) conf=0.77
[21:27:22] ROOT_CAUSE_GENERATED -> The worker-api pod was OOMKilled because its configured memory
limit of 256Mi was insufficient for the application's ope
[21:28:50] SOLUTION_GENERATED -> 4 step(s), 4 command(s)
[21:28:50] RISK_ASSESSED -> safe -- 0 flagged
[21:34:53] WORKFLOW_COMPLETED -> Completed in 461835ms
```