

LogPulse AI - Incident Report

Generated : 2026-06-01 13:25:27 UTC

Trace ID : be173578-b37d-4a74-a788-7d8b26d243b7

1 - Incident Summary

Incident ID : be173578
Category : Memory
Severity : SeverityLevel.FATAL
Source : Kubernetes
Cache : HIT

2 - Raw Log

2026-01-15T22:14:07Z FATAL kubelet: Pod worker-api-7d9f8b-xk2q OOMKilled. Container exceeded memory limit of 256Mi. Exit code 137. Restart count: 14. Node: worker-1. Namespace: production.

3 - Classification

Category : Memory
Severity : SeverityLevel.FATAL
Summary : Kubernetes pod killed due to OOMKilled container exceeding its memory limit.
Source : Kubernetes

4 - Diagnosis

Root Cause : The container's configured memory limit of 256Mi is insufficient for the application's actual memory requirements, leading to repeated OOMKills.
Confidence : 95%
Reasoning : The raw log explicitly states OOMKilled due to exceeding the 256Mi memory limit, which is directly supported by the classification and the highly confident playbook titled 'Pod OOMKilled: Insufficient Memory Limit,' whose primary resolution is to increase the memory limit.
Source Tier : Ilm
Used History : No
Used Playbook : Yes

5 - Proposed Solution

Explanation:

The container's configured memory limit of 256Mi is insufficient for the application's actual memory requirements, leading to repeated OOMKills.

Steps:

- 1. Retrieve the detailed YAML configuration and status of the affected pod to verify the OOMKilled exit code.
- 2. Fetch the logs from the previous failed container instance to analyze memory usage patterns prior to termination.
- 3. Patch the deployment to increase the memory limit to 512Mi and requests to 256Mi, triggering a rolling update.
- 4. Monitor the rollout status of the updated deployment to ensure successful deployment of the new pods.

Commands:

```
$ kubectl get pod worker-api-7d9f8b-xk2q -n production -o yaml
$ kubectl logs worker-api-7d9f8b-xk2q -n production --previous --tail=100
$ kubectl patch deployment worker-api -n production -p
'{"spec":{"template":{"spec":{"containers":[{"name":"worker-api","resources":{"limits":{"memory":"512Mi"},"requests":{"memory":"256Mi"}}}]}}}}'
$ kubectl rollout status deployment/worker-api -n production
```

6 - Risk Assessment

Risk Level : RISKLEVEL.SAFE

Safe Commands (4):

```
[OK] kubectl get pod worker-api-7d9f8b-xk2q -n production -o yaml
[OK] kubectl logs worker-api-7d9f8b-xk2q -n production --previous --tail=100
[OK] kubectl patch deployment worker-api -n production -p
'{"spec":{"template":{"spec":{"containers":[{"name":"worker-api","resources":{"limits":{"memory":"512Mi"},"requests":{"memory":"256Mi"}}}]}}}}'
[OK] kubectl rollout status deployment/worker-api -n production
```

Flagged Commands (0):

(none)

7 - Operator Decision

Operator rejected -- no execution.

8 - Historical Context

Times Seen : 4

First occurrence -- no prior incidents.

9 - Incident Timeline

```
[10:19:29] INCIDENT_RECEIVED -> Incident received (trace_id=be173578...)
[10:19:29] CACHE_HIT -> Cache hit -- seen 4x
[10:19:34] RAG_RETRIEVED -> Playbook retrieved (filtered) -- Memory conf=0.82
[10:19:38] ROOT_CAUSE_GENERATED -> The container's configured memory limit of 256Mi is insufficient for the application's actual memory requirements, leading to OOMKilled
[10:20:27] SOLUTION_GENERATED -> 4 step(s), 4 command(s)
[10:20:27] RISK_ASSESSED -> safe -- 0 flagged
```

[10:25:27] WORKFLOW_COMPLETED -> Completed in 357527ms