

LogPulse AI - Incident Report

Generated : 2026-06-01 16:12:35 UTC

Trace ID : c405bea4-90ed-490b-ada0-a7a0570e5a90

1 - Incident Summary

Incident ID : c405bea4
Category : Unknown
Severity : SeverityLevel.INFO
Source : Fallback
Cache : MISS

2 - Raw Log

2026-01-15T22:14:07Z FATAL kubelet: Pod worker-api-7d9f8b-xk2q OOMKilled. Container exceeded memory limit of 256Mi. Exit code 137. Restart count: 14. Node: worker-1. Namespace: production.

3 - Classification

Category : Unknown
Severity : SeverityLevel.INFO
Summary : Fallback classification triggered due to inference failure.
Source : Fallback

4 - Diagnosis

Root Cause : The 'worker-api' pod's container is repeatedly OOMKilled because its configured memory limit of 256Mi is insufficient for its operational requirements.
Confidence : 90%
Reasoning : The raw log explicitly states the pod was OOMKilled for exceeding its 256Mi memory limit, with a high restart count (14). This directly aligns with the retrieved playbook for 'Pods OOMKilled due to insufficient memory', which prioritizes increasing the memory limit as a resolution.
Source Tier : llm
Used History : No
Used Playbook : Yes

5 - Proposed Solution

Explanation:

The worker-api pod is repeatedly OOMKilled because its configured memory limit of 256Mi is insufficient for its operational requirements.

Steps:

- 1. Inspect the pod details, status, and lifecycle events to confirm OOMKilled status.
- 2. Retrieve logs from the previous crashed instance of the pod to diagnose memory usage prior to termination.
- 3. Increase the memory limit to 1Gi and request to 512Mi for the worker-api deployment to prevent OOM kills.
- 4. Monitor the rollout status of the updated worker-api deployment.

Commands:

```
$ kubectl describe pod worker-api-7d9f8b-xk2q -n production
$ kubectl logs worker-api-7d9f8b-xk2q -n production --previous --tail=100
$ kubectl set resources deployment/worker-api -n production --limits=memory=1Gi
--requests=memory=512Mi
$ kubectl rollout status deployment/worker-api -n production
```

6 - Risk Assessment

Risk Level : RISKLEVEL.SAFE

Safe Commands (4):

- [OK] kubectl describe pod worker-api-7d9f8b-xk2q -n production
- [OK] kubectl logs worker-api-7d9f8b-xk2q -n production --previous --tail=100
- [OK] kubectl set resources deployment/worker-api -n production --limits=memory=1Gi --requests=memory=512Mi
- [OK] kubectl rollout status deployment/worker-api -n production

Flagged Commands (0):

(none)

7 - Operator Decision

Decision : OPERATORDECISION.APPROVED

Executed At : 2026-06-01T13:12:35.686012+00:00

Sandbox Results:

Command : kubectl describe pod worker-api-7d9f8b-xk2q -n production

Exit : 0

Time : 0ms

Output:

Name:	worker-api-7d9f8b-xk2q
Namespace:	production
Node:	worker-1
Status:	OOMKilled
Containers:	
app:	
Image:	app:latest
Limits:	
memory:	256Mi
cpu:	500m
Usage:	
memory:	241Mi

```

    cpu:      245m
  Last State:
    Reason:    OOMKilled
  ... [truncated]

```

Command : `kubectrl logs worker-api-7d9f8b-xk2q -n production --previous --tail=100`

Exit : 0

Time : 0ms

Output:

```

2025-01-15T22:10:00Z WARN Memory usage rising: 78% of 256Mi
2025-01-15T22:12:30Z WARN Memory critical: 91% of 256Mi
2025-01-15T22:13:55Z ERROR Memory critical: 94% of 256Mi -- GC thrashing
2025-01-15T22:14:07Z FATAL OOMKilled: container exceeded 256Mi limit

```

Command : `kubectrl set resources deployment/worker-api -n production --limits=memory=1Gi
--requests=memory=512Mi`

Exit : 0

Time : 0ms

Output:

```

deployment.apps/worker-api resource requirements updated

```

Command : `kubectrl rollout status deployment/worker-api -n production`

Exit : 0

Time : 0ms

Output:

```

Waiting for deployment "worker-api" rollout to finish: 0 of 1 updated replicas are available...

```

8 - Historical Context

Times Seen : 1

First occurrence -- no prior incidents.

9 - Incident Timeline

```

[13:10:38] INCIDENT_RECEIVED -> Incident received (trace_id=c405bea4...)
[13:10:38] CACHE_MISS -> Cache miss -- routing to classifier
[13:10:40] INCIDENT_CLASSIFIED -> Fallback / Info
[13:10:43] RAG_RETRIEVED -> Playbook retrieved (fallback) conf=0.79
[13:10:50] ROOT_CAUSE_GENERATED -> The 'worker-api' pod's container is repeatedly OOMKilled
because its configured memory limit of 256Mi is insufficient fo
[13:12:04] SOLUTION_GENERATED -> 4 step(s), 4 command(s)
[13:12:04] RISK_ASSESSED -> safe -- 0 flagged
[13:12:35] SANDBOX_EXECUTED -> kubectrl describe pod worker-api-7d9f8b-xk2q -n production (exit 0)
[13:12:35] SANDBOX_EXECUTED -> kubectrl logs worker-api-7d9f8b-xk2q -n production --previous
--tail=100 (exit 0)
[13:12:35] SANDBOX_EXECUTED -> kubectrl set resources deployment/worker-api -n production
--limits=memory=1Gi -- (exit 0)
[13:12:35] SANDBOX_EXECUTED -> kubectrl rollout status deployment/worker-api -n production (exit 0)

```

[13:12:35] WORKFLOW_COMPLETED -> Completed in 117591ms