

LogPulse AI - Incident Report

Generated : 2026-06-01 00:40:44 UTC

Trace ID : d6eb1fcc-217e-4a6b-921e-bb0eb39bb8c4

1 - Incident Summary

Incident ID : d6eb1fcc
Category : Unknown
Severity : SeverityLevel.INFO
Source : Fallback
Cache : MISS

2 - Raw Log

2026-05-30 ERROR kubelet: Pod worker-api OOMKilled exceeded 256Mi limit

3 - Classification

Category : Unknown
Severity : SeverityLevel.INFO
Summary : Fallback classification triggered due to inference failure.
Source : Fallback

4 - Diagnosis

Root Cause : The worker-api pod was OOMKilled because its configured memory limit of 256Mi was insufficient for its operational requirements.
Confidence : 95%
Reasoning : The raw log explicitly states 'Pod worker-api OOMKilled exceeded 256Mi limit', and the retrieved playbook is titled 'Pods OOMKilled due to insufficient memory', directly matching the event and suggesting increasing memory limits as a primary resolution.
Source Tier : llm
Used History : No
Used Playbook : Yes

5 - Proposed Solution

Explanation:

The worker-api pod was OOMKilled because its configured memory limit of 256Mi was insufficient for its operational requirements.

Steps:

- 1. Describe the affected pod to inspect its termination state, exit code, and resource limits.
- 2. Retrieve logs from the previous instance of the container prior to the OOMKill event.
- 3. Increase the memory limit of the worker-api deployment to 512Mi and requests to 256Mi to prevent future OOMKills.
- 4. Monitor the rollout status of the updated worker-api deployment.

Commands:

```
$ kubectl describe pod worker-api-7d9f8b-xk2q -n production
$ kubectl logs worker-api-7d9f8b-xk2q -n production --previous
$ kubectl set resources deployment/worker-api -n production --limits=memory=512Mi
--requests=memory=256Mi
$ kubectl rollout status deployment/worker-api -n production
```

6 - Risk Assessment

Risk Level : RISKLEVEL.SAFE

Safe Commands (4):

```
[OK] kubectl describe pod worker-api-7d9f8b-xk2q -n production
[OK] kubectl logs worker-api-7d9f8b-xk2q -n production --previous
[OK] kubectl set resources deployment/worker-api -n production --limits=memory=512Mi --requests=memory=256Mi
[OK] kubectl rollout status deployment/worker-api -n production
```

Flagged Commands (0):

(none)

7 - Operator Decision

Decision : OPERATORDECISION.APPROVED

Executed At : 2026-05-31T21:40:43.547862+00:00

Sandbox Results:

Command : kubectl describe pod worker-api-7d9f8b-xk2q -n production

Exit : 0

Time : 0ms

Output:

```
Name: worker-api-7d9f8b-xk2q
Namespace: production
Node: worker-1
Status: OOMKilled
Containers:
  app:
    Image: app:latest
    Limits:
      memory: 256Mi
      cpu: 500m
    Usage:
      memory: 241Mi
```

```
cpu:      245m
Last State:
Reason:    OOMKilled
... [truncated]
```

Command : kubectl logs worker-api-7d9f8b-xk2q -n production --previous

Exit : 0

Time : 0ms

Output:

```
2025-01-15T22:10:00Z WARN Memory usage rising: 78% of 256Mi
2025-01-15T22:12:30Z WARN Memory critical: 91% of 256Mi
2025-01-15T22:13:55Z ERROR Memory critical: 94% of 256Mi -- GC thrashing
2025-01-15T22:14:07Z FATAL OOMKilled: container exceeded 256Mi limit
```

Command : kubectl set resources deployment/worker-api -n production --limits=memory=512Mi --requests=memory=256Mi

Exit : 0

Time : 0ms

Output:

```
deployment.apps/worker-api resource requirements updated
```

Command : kubectl rollout status deployment/worker-api -n production

Exit : 0

Time : 0ms

Output:

```
Waiting for deployment "worker-api" rollout to finish: 0 of 1 updated replicas are available...
```

8 - Historical Context

Times Seen : 1

First occurrence -- no prior incidents.

9 - Incident Timeline

```
[21:38:39] INCIDENT_RECEIVED -> Incident received (trace_id=d6eb1fcc...)
[21:38:39] CACHE_MISS -> Cache miss -- routing to classifier
[21:38:41] INCIDENT_CLASSIFIED -> Fallback / Info
[21:38:43] RAG_RETRIEVED -> Playbook retrieved (fallback) conf=0.77
[21:38:48] ROOT_CAUSE_GENERATED -> The worker-api pod was OOMKilled because its configured memory limit of 256Mi was insufficient for its operational requi
[21:38:58] SOLUTION_GENERATED -> 4 step(s), 4 command(s)
[21:38:58] RISK_ASSESSED -> safe -- 0 flagged
[21:40:43] SANDBOX_EXECUTED -> kubectl describe pod worker-api-7d9f8b-xk2q -n production (exit 0)
[21:40:43] SANDBOX_EXECUTED -> kubectl logs worker-api-7d9f8b-xk2q -n production --previous (exit 0)
[21:40:43] SANDBOX_EXECUTED -> kubectl set resources deployment/worker-api -n production --limits=memory=512Mi (exit 0)
[21:40:43] SANDBOX_EXECUTED -> kubectl rollout status deployment/worker-api -n production (exit 0)
```

[21:40:43] WORKFLOW_COMPLETED -> Completed in 124063ms