

LogPulse AI - Incident Report

Generated : 2026-06-01 16:09:10 UTC

Trace ID : e2a69026-9d2f-4c31-9aab-d772207fac2e

1 - Incident Summary

Incident ID : e2a69026
Category : Unknown
Severity : SeverityLevel.INFO
Source : Fallback
Cache : MISS

2 - Raw Log

2026-06-01T13:04:18.892Z [k8s-cluster-events] WARNING: Pod worker-api-7f89db5f4c-x2ml9 is consuming 98.4% of its allocated memory limit (256Mi).

3 - Classification

Category : Unknown
Severity : SeverityLevel.INFO
Summary : Fallback classification triggered due to inference failure.
Source : Fallback

4 - Diagnosis

Root Cause : The pod 'worker-api' is consuming 98.4% of its allocated memory limit (256Mi), indicating that the configured memory limit is insufficient for its current operational requirements.
Confidence : 90%
Reasoning : The raw log explicitly states the pod is at 98.4% of its memory limit, directly matching the 'Pods OOMKilled due to insufficient memory' playbook, which prioritizes increasing memory limits.
Source Tier : Ilm
Used History : No
Used Playbook : Yes

5 - Proposed Solution

Explanation:

6 - Risk Assessment

Risk Level : RISKLEVEL.SAFE

Safe Commands (0):

(none)

Flagged Commands (0):

(none)

7 - Operator Decision

Decision : OPERATORDECISION.APPROVED

Executed At : 2026-06-01T13:09:07.854097+00:00

8 - Historical Context

Times Seen : 1

First occurrence -- no prior incidents.

9 - Incident Timeline

[13:07:31] INCIDENT_RECEIVED -> Incident received (trace_id=e2a69026...)
[13:07:31] CACHE_MISS -> Cache miss -- routing to classifier
[13:07:33] INCIDENT_CLASSIFIED -> Fallback / Info
[13:07:47] RAG_RETRIEVED -> Playbook retrieved (fallback) conf=0.72
[13:07:55] ROOT_CAUSE_GENERATED -> The pod 'worker-api' is consuming 98.4% of its allocated memory limit (256Mi), indicating that the configured memory lim
[13:08:44] SOLUTION_GENERATED -> 0 step(s), 0 command(s)
[13:08:44] RISK_ASSESSED -> safe -- 0 flagged
[13:09:07] WORKFLOW_COMPLETED -> Completed in 96129ms