

LogPulse AI - Incident Report

Generated : 2026-05-29 17:11:24 UTC

Trace ID : ed74d43f-4843-4be8-b5dc-1ed7afa22daa

1 - Incident Summary

Incident ID : ed74d43f
Category : Memory
Severity : SeverityLevel.FATAL
Source : Kubernetes
Cache : MISS

2 - Raw Log

17:10:37 FATAL OOMKilled: pod/worker-api exceeded memory limit 256Mi restarts=14

3 - Classification

Category : Memory
Severity : SeverityLevel.FATAL
Summary : Kubernetes worker API pod exceeded its memory limit and was terminated.
Source : Kubernetes

4 - Diagnosis

Root Cause : The pod/worker-api is repeatedly OOMKilled because its configured memory limit of 256Mi is insufficient for the application's runtime memory requirements, leading to 14 restarts.
Confidence : 98%
Reasoning : The raw log explicitly states 'OOMKilled: pod/worker-api exceeded memory limit 256Mi restarts=14'. The retrieved playbook directly addresses 'Insufficient Memory Limit', and the incident history confirms this exact root cause for previous occurrences of this incident.
Source Tier : llm
Used History : Yes
Used Playbook : Yes

5 - Proposed Solution

Explanation:

Gemini unavailable; remediation derived from playbook 'Pod OOMKilled: Insufficient Memory Limit' (RAG confidence=0.8064). Review the playbook steps and adapt commands to the live environment before execution.

Steps:

- 1. Verify the `OOMKilled` status using `kubectl get pod <pod-name> -o yaml`.
- 2. Check container memory usage metrics (e.g., Prometheus/Grafana) to identify peak memory consumption.
- 3. Analyze application logs for clues about memory spikes or growth patterns.
- 4. Increase the `resources.limits.memory` for the affected container in its Deployment/StatefulSet manifest.
- 5. Roll out the updated deployment using `kubectl apply -f <manifest-file>` or `kubectl rollout restart deployment/<deployment-name>`.
- 6. name: my-app

Commands:

```
$ kubectl get pod worker-api-7d9f8b-xk2q -o yaml
$ kubectl apply -f <manifest-file>
$ kubectl rollout restart deployment/worker-api
```

6 - Risk Assessment

Risk Level : **RISKLEVEL.WARNING**

Safe Commands (2):

[OK] kubectl get pod worker-api-7d9f8b-xk2q -o yaml
[OK] kubectl apply -f <manifest-file>

Flagged Commands (1):

[!] kubectl rollout restart deployment/worker-api [kubectl\s+rollout\s+restart]

7 - Operator Decision

Decision : OPERATORDECISION.APPROVED

Executed At : 2026-05-29T14:11:24.941543+00:00

Sandbox Results:

Command : kubectl get pod worker-api-7d9f8b-xk2q -o yaml
Exit : 0
Time : 0ms

Output:

```
[auto-approve] simulated: kubectl get pod worker-api-7d9f8b-xk2q -o yaml
```

Command : kubectl apply -f <manifest-file>
Exit : 0
Time : 0ms

Output:

```
[auto-approve] simulated: kubectl apply -f <manifest-file>
```

Command : kubectl rollout restart deployment/worker-api
Exit : 0
Time : 0ms

Output:

```
[auto-approve] simulated: kubectl rollout restart deployment/worker-api
```

8 - Historical Context

Times Seen : 1

Past Outcomes:

- analyzed on 2026-05-29 13:56:38.610496+00:00

9 - Incident Timeline

[14:11:24] WORKFLOW_COMPLETED -> Completed in 20477ms