

LogPulse AI - Incident Report

Generated : 2026-06-01 11:13:42 UTC

Trace ID : f803e2d1-43a9-4a41-a6eb-4aff6acbc404

1 - Incident Summary

Incident ID : f803e2d1
Category : Unknown
Severity : SeverityLevel.INFO
Source : Fallback
Cache : MISS

2 - Raw Log

2026-05-30 18:22:11 ERROR kubelet: Pod worker-api-7d9f8b-xk2q in namespace production OOMKilled (container=worker-api) exceeded 256Mi limit. Restart count: 14. Memory pressure detected on node worker-1.

3 - Classification

Category : Unknown
Severity : SeverityLevel.INFO
Summary : Fallback classification triggered due to inference failure.
Source : Fallback

4 - Diagnosis

Root Cause : The worker-api pod is being OOMKilled because its configured memory limit of 256Mi is insufficient for its operational requirements, leading to repeated restarts.

Confidence : 95%

Reasoning : The raw log explicitly states the pod was OOMKilled after exceeding its 256Mi memory limit and has a high restart count, which directly matches the retrieved playbook for 'Pods OOMKilled due to insufficient memory' and its suggested resolution of increasing the memory limit.

Source Tier : llm
Used History : No
Used Playbook : Yes

5 - Proposed Solution

Explanation:

Gemini unavailable; remediation derived from playbook 'Pods OOMKilled due to insufficient memory' (RAG

confidence=0.7871). Review the playbook steps and adapt commands to the live environment before execution.

Steps:

- 1. Identify the specific pods and deployments experiencing OOMKills using `kubectl get events` and `kubectl describe pod <pod-name>`.
- 2. Check the application logs within the affected pods for memory-related errors.
- 3. Review the `memory.limit` and `memory.request` defined in the pod's container spec.
- 4. If the application is known to have increased memory requirements, increment the `memory.limit` for the affected deployment in its YAML.
- 5. If a memory leak is suspected, collaborate with the development team to analyze heap dumps or application-specific memory metrics.
- 6. Consider scaling out the deployment or the node pool if node-wide memory pressure is observed consistently.
- 7. name: my-container

Commands:

```
$ kubectl apply -f deployment.yaml
$ kubectl rollout status deployment/my-service -w
```

6 - Risk Assessment

Risk Level : **RISKLEVEL.SAFE**

Safe Commands (2):

- [OK] `kubectl apply -f deployment.yaml`
- [OK] `kubectl rollout status deployment/my-service -w`

Flagged Commands (0):

(none)

7 - Operator Decision

Operator rejected -- no execution.

8 - Historical Context

Times Seen : 1
First occurrence -- no prior incidents.

9 - Incident Timeline

```
[08:13:40] WORKFLOW_COMPLETED -> Completed in 591109ms
```